

# 1.Problem Analysis

In the estate market, each house has a selling price based on its different features and conditions. Based on the given train dataset, test dataset and description of each features, this project aims to predict the provided house sale price on the market.

## 2.Get Data and Data cleaning

Import the pandas library and load train.csv file into Python. A quick review of the input data is as follows: there are 81 different features in column and 1460 observations samples in row:

```
>>> df_train
   Id  MSSubClass MSZoning  LotFrontage  LotArea  Street  Alley  LotShape  LandContour  Utilities  ... PoolArea PoolQC  Fence MiscFeature MiscVal MoSold YrSold  SaleType  SaleCondition  SalePrice
0    1         60      RL         65.0     8450   Pave   NaN    Reg         Lvl     AllPub  ...      0   NaN   NaN      NaN      0      2    2008      WD      Normal      20850
1    2         20      RL         80.0     9600   Pave   NaN    Reg         Lvl     AllPub  ...      0   NaN   NaN      NaN      0      5    2007      WD      Normal      18150
2    3         60      RL         68.0    11250   Pave   NaN  IR1         Lvl     AllPub  ...      0   NaN   NaN      NaN      0      9    2008      WD      Normal      22350
3    4         70      RL         60.0     9550   Pave   NaN  IR1         Lvl     AllPub  ...      0   NaN   NaN      NaN      0      2    2006      WD      Abnormal      14000
4    5         60      RL         84.0    14260   Pave   NaN  IR1         Lvl     AllPub  ...      0   NaN   NaN      NaN      0     12    2008      WD      Normal      25000
...  ...      ...      ...      ...      ...      ...      ...      ...      ...      ...  ...  ...  ...      ...      ...      ...      ...      ...      ...
1455 1455         60      RL         62.0     7917   Pave   NaN    Reg         Lvl     AllPub  ...      0   NaN   NaN      NaN      0      8    2007      WD      Normal      17500
1456 1456         20      RL         85.0    13175   Pave   NaN    Reg         Lvl     AllPub  ...      0   NaN  MnPrv      NaN      0      2    2010      WD      Normal      21000
1457 1457         70      RL         66.0     9042   Pave   NaN    Reg         Lvl     AllPub  ...      0   NaN  GdPrv     Shed    2500      5    2010      WD      Normal      26650
1458 1458         20      RL         68.0     9717   Pave   NaN    Reg         Lvl     AllPub  ...      0   NaN   NaN      NaN      0      4    2010      WD      Normal      14212
1459 1459         20      RL         75.0     9937   Pave   NaN    Reg         Lvl     AllPub  ...      0   NaN   NaN      NaN      0      6    2008      WD      Normal      14750


[1460 rows x 81 columns]
```

All column headers in the train dataset can be observed as below:

```
>>> df_train=pd.read_csv('train.csv')
>>> print(df_train.columns)
Index(['Id', 'MSSubClass', 'MSZoning', 'LotFrontage', 'LotArea', 'Street',
       'Alley', 'LotShape', 'LandContour', 'Utilities', 'LotConfig',
       'LandSlope', 'Neighborhood', 'Condition1', 'Condition2', 'BldgType',
       'HouseStyle', 'OverallQual', 'OverallCond', 'YearBuilt', 'YearRemodAdd',
       'RoofStyle', 'RoofMatl', 'Exterior1st', 'Exterior2nd', 'MasVnrType',
       'MasVnrArea', 'ExterQual', 'ExterCond', 'Foundation', 'BsmtQual',
       'BsmtCond', 'BsmtExposure', 'BsmtFinType1', 'BsmtFinSF1',
       'BsmtFinType2', 'BsmtFinSF2', 'BsmtUnfSF', 'TotalBsmtSF', 'Heating',
       'HeatingQC', 'CentralAir', 'Electrical', '1stFlrSF', '2ndFlrSF',
       'LowQualFinSF', 'GrLivArea', 'BsmtFullBath', 'BsmtHalfBath', 'FullBath',
       'HalfBath', 'BedroomAbvGr', 'KitchenAbvGr', 'KitchenQual',
       'TotRmsAbvGrd', 'Functional', 'Fireplaces', 'FireplaceQu', 'GarageType',
       'GarageYrBlt', 'GarageFinish', 'GarageCars', 'GarageArea', 'GarageQual',
       'GarageCond', 'PavedDrive', 'WoodDeckSF', 'OpenPorchSF',
       'EnclosedPorch', '3SsnPorch', 'ScreenPorch', 'PoolArea', 'PoolQC',
       'Fence', 'MiscFeature', 'MiscVal', 'MoSold', 'YrSold', 'SaleType',
       'SaleCondition', 'SalePrice'],
      dtype='object')
>>>
```


A descriptive statistic summary of SalePrice and distribution plot in histogram are as below:

```
>>> df_train['SalePrice'].describe()
count      1460.000000
mean     180921.195890
std       79442.502883
min       34900.000000
25%      129975.000000
50%      163000.000000
75%      214000.000000
max       755000.000000
Name: SalePrice, dtype: float64
>>>
```




From the statistics and SalePrice distribution perspective, the mean and max value of SalePrice is about \$180921 and \$755000, and most of the SalePrice values are between \$100000 to \$400000.

Through correlation matrix, the correlation between each two variables can be observed:



Based on the above Heatmap figure, a low correlation can be observed between feature EnclosedPorch and SalePrice. A further scatter plot is applied as below which further shows that there is no clear linear correlation between this two feature. EnclosedPorch will be kept until further confirmation.

This strategy can be applied to check whether a feature is highly related with SalePrice or not.



Duplicated observation in sample dataset can be showed by using duplicated() function. The output will be Empty DataFrame when there is no duplicated rows in the dataset.

```
>>> df_train[df_train.duplicated()==True]
Empty DataFrame
Columns: [Id, MSSubClass, MSZoning, LotFrontage, LotArea, Street, Alley, LotShape, LandContour, Utilities, LotConfig, LandSlope, Neighborhood, Condition1, Condition2, BldgType, HouseStyle, OverallQual, OverallCond, YearBuilt, YearRemodAdd, RoofStyle, RoofMatl, Exterior1st, Exterior2nd, MasVnrType, MasVnrArea, ExterQual, ExterCond, Foundation, BsmtQual, BsmtCond, BsmtExposure, BsmtFinType1, BsmtFinSF1, BsmtFinType2, BsmtFinSF2, BsmtUnfSF, TotalBsmtSF, Heating, HeatingQC, CentralAir, Electrical, 1stFlrSF, 2ndFlrSF, LowQualFinSF, GrLivArea, BsmtFullBath, BsmtHalfBath, FullBath, HalfBath, BedroomAbvGr, KitchenAbvGr, KitchenQual, TotRmsAbvGrd, Functional, Fireplaces, FireplaceQu, GarageType, GarageYrBlt, GarageFinish, GarageCars, GarageArea, GarageQual, GarageCond, PavedDrive, WoodDeckSF, OpenPorchSF, EnclosedPorch, 3SsnPorch, ScreenPorch, PoolArea, PoolQC, Fence, MiscFeature, MiscVal, MoSold, YrSold, SaleType, SaleCondition, SalePrice]
Index: []
>>>
```

There are mainly 3 data types in the train dataset, which are int64, object and float64.

The detailed info can be observed as below snapshots:

```
>>> res = df_train.dtypes
>>> print(res[res == np.dtype('int64')])
Id                int64
MSSubClass        int64
LotArea           int64
OverallQual       int64
OverallCond       int64
YearBuilt         int64
YearRemodAdd      int64
BsmtFinSF1        int64
BsmtFinSF2        int64
BsmtUnfSF         int64
TotalBsmtSF       int64
1stFlrSF          int64
2ndFlrSF          int64
LowQualFinSF      int64
GrLivArea         int64
BsmtFullBath      int64
BsmtHalfBath      int64
FullBath          int64
HalfBath          int64
BedroomAbvGr      int64
KitchenAbvGr      int64
TotRmsAbvGrd      int64
Fireplaces        int64
GarageCars        int64
GarageArea        int64
WoodDeckSF        int64
OpenPorchSF       int64
EnclosedPorch     int64
3SsnPorch         int64
ScreenPorch       int64
PoolArea          int64
MiscVal           int64
MoSold            int64
YrSold            int64
SalePrice         int64
dtype: object
>>>
```


```
>>> print(res[res == np.dtype('float64')])
LotFrontage       float64
MasVnrArea        float64
GarageYrBlt       float64
dtype: object
>>>
```

```
>>> print(res[res==np.dtype('object')])
MSZoning          object
Street            object
Alley             object
LotShape          object
LandContour       object
Utilities         object
LotConfig         object
LandSlope         object
Neighborhood      object
Condition1        object
Condition2        object
BldgType          object
HouseStyle        object
RoofStyle         object
RoofMatl          object
Exterior1st       object
Exterior2nd       object
MasVnrType        object
ExterQual         object
ExterCond         object
Foundation        object
BsmtQual          object
BsmtCond          object
BsmtExposure      object
BsmtFinType1      object
BsmtFinType2      object
Heating           object
HeatingQC         object
CentralAir        object
Electrical        object
KitchenQual       object
Functional        object
FireplaceQu       object
GarageType        object
GarageFinish      object
GarageQual        object
GarageCond        object
PavedDrive        object
PoolQC           object
Fence             object
MiscFeature       object
SaleType          object
SaleCondition     object
dtype: object
>>>
```



The unique values of each variable can also be observed through unique() function as below example:

```
>>> print(df_train["LotConfig"].unique())
['Inside' 'FR2' 'Corner' 'CulDSac' 'FR3']
>>>
```

For numerical variable, feature scaling can be applied to fit and transform data into normal distribution .



Data transformation: the distribution of GrLivArea column can be observed with fitting normalisation format as the left figure. GrLivArea column can also be calculated through log() function and a new distribution picture can be plotted as the right figure.



About missing data processing:

Sum all null values of each column and calculate the percentage of missing values for each column, then the first 20 highest percentages of

missing values for columns can be found out and further processing can be applied based on different situations, for example, for any column whose missing percentage is greater than 15% will be dropped.

```
>>> total = df_train.isnull().sum().sort_values(ascending=False)
>>> total
PoolQC      1453
MiscFeature 1406
Alley       1369
Fence       1179
FireplaceQu 690
...
CentralAir  0
SaleCondition 0
Heating     0
TotalBsmtSF 0
Id          0
Length: 81, dtype: int64
>>>
```

```
>>> missing_data = pd.concat([total, percent], axis=1, keys=['Total', 'Percent'])
>>> missing_data.head(20)
   Total  Percent
PoolQC   1453  0.995205
MiscFeature 1406  0.963014
Alley    1369  0.937671
Fence    1179  0.807534
FireplaceQu 690  0.472603
LotFrontage 259  0.177397
GarageCond  81  0.055479
GarageType  81  0.055479
GarageYrBlt 81  0.055479
GarageFinish 81  0.055479
GarageQual  81  0.055479
BsmtExposure 38  0.026027
BsmtFinType2 38  0.026027
BsmtFinType1 37  0.025342
BsmtCond    37  0.025342
BsmtQual    37  0.025342
MasVnrArea   8  0.005479
MasVnrType   8  0.005479
Electrical   1  0.000685
Utilities    0  0.000000
>>>
```

```
>>> percent = (df_train.isnull().sum()/df_train.isnull().count()).sort_values(ascending=False)
>>> percent
PoolQC      0.995205
MiscFeature  0.963014
Alley        0.937671
Fence        0.807534
FireplaceQu  0.472603
...
CentralAir   0.000000
SaleCondition 0.000000
Heating      0.000000
TotalBsmtSF  0.000000
Id           0.000000
Length: 81, dtype: float64
>>>
```

Missing data can also be imputed: The method is that set the values of a highly correlated column to replace the missing values.

There are 259 values are missing before the imputation, and after processing the missing value number is 0.

```
>>>
>>> print(df_train["LotFrontage"].isnull().sum())
259
>>> cond = df_train['LotFrontage'].isnull()
>>> df_train["LotFrontage"][cond]=df_train["SqrtLotArea"][cond]
__main__:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
>>> print(df_train["LotFrontage"].isnull().sum())
0
>>> █
```

Another way is to mark the missing value of the column as “Missing”:

```
>>> mis=df_train['GarageType'].isnull()
>>>
>>> df_train["GarageType"][mis]="Missing"
>>> df_train["GarageType"].unique()
array(['Attchd', 'Detchd', 'BuiltIn', 'CarPort', 'Missing', 'Basement',
       '2Types'], dtype=object)
>>> █
```

Drop the columns whose data value missing percentage is greater than 15%, drop the row whose “Electrical” is null. After dropping, the data frame is changed into 1459 rows \* 76 columns from original 1460 rows \* 81 columns.

```
>>> df_train = df_train.drop((missing_data[missing_data['Percent'] > 0.15]).index,1)
>>> df_train = df_train.drop(df_train.loc[df_train['Electrical'].isnull()].index)
>>> df_train
   Id  MSSubClass  MSZoning  LotArea  Street  LotShape  LandContour  Utilities  LotConfig  ...  ScreenPorch  PoolArea  MiscVal  MoSold  YrSold  SaleType  SaleCondition  SalePrice  SqrtLotArea
0    1         60      RL      8450   Pave      Reg      Lvl      AllPub   Inside  ...         0         0         0         2    2008      WD      Normal      208500      91.923882
1    2         20      RL      9600   Pave      Reg      Lvl      AllPub   Inside  ...         0         0         0         5    2007      WD      Normal      181500      97.979590
2    3         60      RL     11250   Pave      IR1      Lvl      AllPub   Inside  ...         0         0         0         9    2008      WD      Normal      223500     106.066017
3    4         70      RL      9550   Pave      IR1      Lvl      AllPub   Corner  ...         0         0         0         2    2006      WD      Abnorml      140000      97.724101
4    5         60      RL     14260   Pave      IR1      Lvl      AllPub   FR2     ...         0         0         0        12    2008      WD      Normal      250000     119.415242
...  ...      ...      ...      ...      ...      ...      ...      ...      ...  ...      ...      ...      ...      ...      ...      ...      ...      ...
1455 1456         60      RL      7917   Pave      Reg      Lvl      AllPub   Inside  ...         0         0         0         8    2007      WD      Normal      175000      88.977525
1456 1457         20      RL     13175   Pave      Reg      Lvl      AllPub   Inside  ...         0         0         0         2    2010      WD      Normal      210000     114.782483
1457 1458         70      RL      9042   Pave      Reg      Lvl      AllPub   Inside  ...         0         0     2500         5    2010      WD      Normal      266500      95.089432
1458 1459         20      RL      9717   Pave      Reg      Lvl      AllPub   Inside  ...         0         0         0         4    2010      WD      Normal      142125      98.574845
1459 1460         20      RL      9937   Pave      Reg      Lvl      AllPub   Inside  ...         0         0         0         6    2008      WD      Normal      147500      99.684502

[1459 rows x 76 columns]
```

Check correlation between two variables through corr() function: LotFrontage and LotArea

```
>>> df_train['LotFrontage'].corr(df_train['LotArea'])
0.42609501877180833
>>> █
```

Another sqrt() function can be further applied to reduce the numeric effect.

```
>>> df_train['SqrtLotArea']=np.sqrt(df_train['LotArea'])
>>> df_train['LotFrontage'].corr(df_train['SqrtLotArea'])
0.6020022167939361
>>> █
```

For outliers processing:

There are some outliers for several columns, these outliers can be removed directly to reduce the effects of outliers in final prediction processing. The train data is changed into 1444 \* 76 from 1459 \* 76.

